

- S. Perkowitz. The computers take over. In *Hollywood Science: Movies, Science, and the End of the World*, pages 142–164. Columbia University Press, 2007.
- A. Radford, J. Wu, R. Child, D. Luan, D. Amodei, and I. Sutskever. Language models are unsupervised multitask learners, 2019. https://cdn.openai.com/better-language-models/language_models_are_unsupervised_multitask_learners.pdf.
- J. W. Rae, S. Borgeaud, T. Cai, K. Millican, J. Hoffmann, et al. Scaling language models: Methods, analysis & insights from training Gopher. *arXiv preprint*, arXiv:2112.11446, 2021.
- L. Reynolds and K. McDonell. Multiversal views on language models. *arXiv preprint*, arXiv:2102.06391, 2021.
- K. Roose. Bing’s A.I. Chat: ‘I Want to Be Alive.’. New York Times, 26th February, 2023. <https://www.nytimes.com/2023/02/16/technology/bing-chatbot-transcript.html>.
- E. Ruane, A. Birhane, and A. Ventresque. Conversational AI: Social and ethical considerations. In *Proceedings 27th AIAI Irish Conference on Artificial Intelligence and Cognitive Science*, pages 104–115, 2019.
- S. Russell and P. Norvig. *Artificial Intelligence: A Modern Approach, 3rd Edition*. Prentice Hall, 2010.
- T. Schick, J. Dwivedi-Yu, R. Dessì, R. Raileanu, M. Lomeli, et al. Toolformer: Language models can teach themselves to use tools. *arXiv preprint arXiv:2302.04761*, 2023.
- M. Shanahan. Talking about large language models. *arXiv preprint*, arXiv:2212.03551, 2023.
- N. Stiennon, L. Ouyang, J. Wu, D. M. Ziegler, R. Lowe, et al. Learning to summarize from human feedback. In *Advances in Neural Information Processing Systems*, pages 3008–3021, 2020.
- R. Thoppilan, D. De Freitas, J. Hall, N. Shazeer, A. Kulshreshtha, et al. LaMDA: Language models for dialog applications. *arXiv preprint*, arXiv:2201.08239, 2022.
- A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, pages 5998–6008, 2017.
- J. Wei, Y. Tay, R. Bommasani, C. Raffel, B. Zoph, et al. Emergent abilities of large language models. *Transactions on Machine Learning Research*, 2022.
- S. Willison. Bing: “I will not harm you unless you harm me first”. Simon Willison’s Weblog, 15th February, 2023. <https://simonwillison.net/2023/Feb/15/bing/>.
- S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, et al. React: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations*, 2023.